

Task4

November 19, 2025

1 Task 4

```
[34]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
import joblib
```

```
[35]: df = pd.read_csv("Telco-Customer-Churn.csv")
df.head()
```

```
[35]:    customerID  gender  SeniorCitizen  Partner  Dependents  tenure  PhoneService \
0  7590-VHVEG  Female           0      Yes        No         1        No
1  5575-GNVDE    Male           0       No        No        34      Yes
2  3668-QPYBK    Male           0       No        No         2      Yes
3  7795-CFOCW    Male           0       No        No        45        No
4  9237-HQITU  Female           0       No        No         2      Yes

      MultipleLines  InternetService  OnlineSecurity  ...  DeviceProtection \
0  No phone service            DSL           No  ...
1          No                 DSL           Yes  ...
2          No                 DSL           Yes  ...
3  No phone service            DSL           Yes  ...
4          No  Fiber optic           No  ...

      TechSupport  StreamingTV  StreamingMovies  ...  Contract  PaperlessBilling \
0          No        No           No  ...  Month-to-month        Yes
1          No        No           No  ...  One year          No
2          No        No           No  ...  Month-to-month        Yes
```

3	Yes	No	No	One year	No	
4	No	No	No	Month-to-month	Yes	
			PaymentMethod	MonthlyCharges	TotalCharges	Churn
0		Electronic check		29.85	29.85	No
1		Mailed check		56.95	1889.5	No
2		Mailed check		53.85	108.15	Yes
3	Bank transfer (automatic)			42.30	1840.75	No
4		Electronic check		70.70	151.65	Yes

[5 rows x 21 columns]


```
[36]: print(df.shape)
df.info()
df.describe()
```

```
(7043, 21)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customerID      7043 non-null   object 
 1   gender          7043 non-null   object 
 2   SeniorCitizen   7043 non-null   int64  
 3   Partner         7043 non-null   object 
 4   Dependents     7043 non-null   object 
 5   tenure          7043 non-null   int64  
 6   PhoneService    7043 non-null   object 
 7   MultipleLines   7043 non-null   object 
 8   InternetService 7043 non-null   object 
 9   OnlineSecurity  7043 non-null   object 
 10  OnlineBackup    7043 non-null   object 
 11  DeviceProtection 7043 non-null   object 
 12  TechSupport     7043 non-null   object 
 13  StreamingTV     7043 non-null   object 
 14  StreamingMovies  7043 non-null   object 
 15  Contract        7043 non-null   object 
 16  PaperlessBilling 7043 non-null   object 
 17  PaymentMethod   7043 non-null   object 
 18  MonthlyCharges  7043 non-null   float64
 19  TotalCharges    7043 non-null   object 
 20  Churn           7043 non-null   object 

dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
[36]: SeniorCitizen      tenure  MonthlyCharges
count    7043.000000  7043.000000    7043.000000
mean     0.162147    32.371149    64.761692
std      0.368612    24.559481    30.090047
min      0.000000    0.000000    18.250000
25%     0.000000    9.000000    35.500000
50%     0.000000   29.000000    70.350000
75%     0.000000   55.000000   89.850000
max     1.000000   72.000000   118.750000
```

```
[37]: df["Churn"].value_counts().plot(kind="bar", color=['blue','red'])
plt.title("Churn Distribution")
plt.show()
```



```
[38]: df = df.replace(" ", np.nan)
df = df.dropna()
df.isnull().sum()
```

```
[38]: customerID      0  
gender          0  
SeniorCitizen   0  
Partner         0  
Dependents     0  
tenure          0  
PhoneService    0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport     0  
StreamingTV    0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
TotalCharges    0  
Churn           0  
dtype: int64
```

```
[39]: df["TotalCharges"] = pd.to_numeric(df["TotalCharges"])
```

```
X = df.drop("Churn", axis=1)  
y = df["Churn"]
```

```
le = LabelEncoder()  
y = le.fit_transform(y)
```

```
[40]: categorical_cols = X.select_dtypes(include=["object"]).columns  
numeric_cols = X.select_dtypes(include=["int64", "float64"]).columns  
categorical_cols, numeric_cols
```

```
[40]: (Index(['customerID', 'gender', 'Partner', 'Dependents', 'PhoneService',  
           'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup',  
           'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies',  
           'Contract', 'PaperlessBilling', 'PaymentMethod'],  
           dtype='object'),  
       Index(['SeniorCitizen', 'tenure', 'MonthlyCharges', 'TotalCharges'],  
           dtype='object'))
```

```
[41]: preprocess = ColumnTransformer(  
    transformers=[  
        ("cat", OneHotEncoder(handle_unknown='ignore'), categorical_cols),  
        ("num", StandardScaler(), numeric_cols)
```

```
    ]  
)
```

```
[42]: X_train, X_test, y_train, y_test = train_test_split(  
        X, y, test_size=0.2, random_state=42  
)  
model = Pipeline(  
    steps=[  
        ("preprocess", preprocess),  
        ("classifier", RandomForestClassifier(n_estimators=200,  
        random_state=42))  
    ]  
)
```


```
[43]: model.fit(X_train, y_train)  
y_pred = model.predict(X_test)
```

```
[44]: accuracy = accuracy_score(y_test, y_pred)  
accuracy
```

```
[44]: 0.7896233120113717
```

```
[45]: print(classification_report(y_test, y_pred))  
cm = confusion_matrix(y_test, y_pred)  
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')  
plt.title("Confusion Matrix")  
plt.show()
```

	precision	recall	f1-score	support
0	0.82	0.91	0.86	1033
1	0.65	0.45	0.53	374
accuracy			0.79	1407
macro avg	0.74	0.68	0.70	1407
weighted avg	0.78	0.79	0.78	1407




```
[46]: rf = model.named_steps["classifier"]
ohe = model.named_steps["preprocess"].named_transformers_["cat"]
encoded_cat_cols = ohe.get_feature_names_out(categorical_cols)

all_features = np.concatenate([encoded_cat_cols, numeric_cols])

importances = rf.feature_importances_

imp_df = pd.DataFrame({
    "Feature": all_features,
    "Importance": importances
}).sort_values(by="Importance", ascending=False)

imp_df.head(20)
imp_df.head(20).plot(kind='bar', x='Feature', y='Importance', figsize=(12,5))
plt.title("Top 20 Important Features")
plt.show()
```



```
[47]: joblib.dump(model, "churn_model.pkl")
```

```
[47]: ['churn_model.pkl']
```